K Nearest Neighbors with Python

You've been given a classified data set from a company! They've hidden the feature column names but have given you the data and the target classes.

We'll try to use KNN to create a model that directly predicts a class for a new data point based off of the features.

Let's grab it and use it!

Import Libraries

```
In [1]: import pandas as pd
        import seaborn as sns
        import matplotlib.pyplot as plt
        import numpy as np
        %matplotlib inline
```

Get the Data

Set index_col=0 to use the first column as the index.

```
In [74]: df = pd.read_csv("Classified Data",index_col=0)
```

```
In [75]: df.head()
```

Out[75]:

	WTT	PTI	EQW	SBI	LQE	QWG	FDJ	PJF	HQE	N
0	0.913917	1.162073	0.567946	0.755464	0.780862	0.352608	0.759697	0.643798	0.879422	1.
1	0.635632	1.003722	0.535342	0.825645	0.924109	0.648450	0.675334	1.013546	0.621552	1.
2	0.721360	1.201493	0.921990	0.855595	1.526629	0.720781	1.626351	1.154483	0.957877	1.
3	1.234204	1.386726	0.653046	0.825624	1.142504	0.875128	1.409708	1.380003	1.522692	1.
4	1.279491	0.949750	0.627280	0.668976	1.232537	0.703727	1.115596	0.646691	1.463812	1.

Standardize the Variables

Because the KNN classifier predicts the class of a given test observation by identifying the observations that are nearest to it, the scale of the variables matters. Any variables that are on a large scale will have a much larger effect on the distance between the observations, and hence on the KNN classifier, than variables that are on a small scale.

```
In [78]: from sklearn.preprocessing import StandardScaler
```

```
In [79]: scaler = StandardScaler()
```

```
In [80]:  scaler.fit(df.drop('TARGET CLASS',axis=1))

Out[80]:  StandardScaler(copy=True, with_mean=True, with_std=True)

In [81]:  scaled_features = scaler.transform(df.drop('TARGET CLASS',axis=1))

In [82]:  df_feat = pd.DataFrame(scaled_features,columns=df.columns[:-1])
          df_feat.head()
```

Out[82]:

	WTT	PTI	EQW	SBI	LQE	QWG	FDJ	PJF	HQE
0	-0.123542	0.185907	-0.913431	0.319629	-1.033637	-2.308375	-0.798951	-1.482368	-0.949
1	-1.084836	-0.430348	-1.025313	0.625388	-0.444847	-1.152706	-1.129797	-0.202240	-1.828
2	-0.788702	0.339318	0.301511	0.755873	2.031693	-0.870156	2.599818	0.285707	-0.682
3	0.982841	1.060193	-0.621399	0.625299	0.452820	-0.267220	1.750208	1.066491	1.2413
4	1.139275	-0.640392	-0.709819	-0.057175	0.822886	-0.936773	0.596782	-1.472352	1.0407

Train Test Split

```
In [83]:  from sklearn.model_selection import train_test_split

In [84]:  X_train, X_test, y_train, y_test = train_test_split(scaled_features,df['TARGET CLA
          SS'],
                                                            test_size=0.30)
```

Using KNN

Remember that we are trying to come up with a model to predict whether someone will TARGET CLASS or not. We'll
start with k=1.

```
In [85]:  from sklearn.neighbors import KNeighborsClassifier

In [86]:  knn = KNeighborsClassifier(n_neighbors=1)

In [87]:  knn.fit(X_train,y_train)

Out[87]:  KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                      metric_params=None, n_jobs=1, n_neighbors=1, p=2,
                      weights='uniform')

In [88]:  pred = knn.predict(X_test)
```

Predictions and Evaluations

Let's evaluate our KNN model!

```
In [89]:  from sklearn.metrics import classification_report,confusion_matrix
```

```
In [90]: print(confusion_matrix(y_test,pred))

         [[125  18]
          [ 13 144]]
```

```
In [91]: print(classification_report(y_test,pred))
```

```
                   precision    recall  f1-score   support

               0        0.91      0.87      0.89       143
               1        0.89      0.92      0.90       157

      avg / total        0.90      0.90      0.90       300
```

Choosing a K Value

Let's go ahead and use the elbow method to pick a good K Value:

```python
In [98]: error_rate = []

         # Will take some time
         for i in range(1,40):

             knn = KNeighborsClassifier(n_neighbors=i)
             knn.fit(X_train,y_train)
             pred_i = knn.predict(X_test)
             error_rate.append(np.mean(pred_i != y_test))
```

```python
In [99]: plt.figure(figsize=(10,6))
         plt.plot(range(1,40),error_rate,color='blue', linestyle='dashed', marker='o',
                  markerfacecolor='red', markersize=10)
         plt.title('Error Rate vs. K Value')
         plt.xlabel('K')
         plt.ylabel('Error Rate')
```

```
Out[99]: <matplotlib.text.Text at 0x11ca82ba8>
```

Here we can see that that after arouns K>23 the error rate just tends to hover around 0.06-0.05 Let's retrain the model with that and check the classification report!

In [100]:
```
# FIRST A QUICK COMPARISON TO OUR ORIGINAL K=1
knn = KNeighborsClassifier(n_neighbors=1)

knn.fit(X_train,y_train)
pred = knn.predict(X_test)

print('WITH K=1')
print('\n')
print(confusion_matrix(y_test,pred))
print('\n')
print(classification_report(y_test,pred))
```

WITH K=1

[[125 18]
 [13 144]]

	precision	recall	f1-score	support
0	0.91	0.87	0.89	143
1	0.89	0.92	0.90	157
avg / total	0.90	0.90	0.90	300

In [101]:
```
# NOW WITH K=23
knn = KNeighborsClassifier(n_neighbors=23)

knn.fit(X_train,y_train)
pred = knn.predict(X_test)

print('WITH K=23')
print('\n')
print(confusion_matrix(y_test,pred))
print('\n')
print(classification_report(y_test,pred))
```

WITH K=23

[[132 11]
 [5 152]]

	precision	recall	f1-score	support
0	0.96	0.92	0.94	143
1	0.93	0.97	0.95	157
avg / total	0.95	0.95	0.95	300

In [ ]: